

**Course #2:**

**Deep Learning, from  
MLP to CNN**

# Roadmap


- Recap from course #1
- MLP and Image classification as a case study
- CNN: basic principles
- Application to image classification
- Classic CNN architectures
- Auto-encoders

# Recap from Course #1

## Things to know

- Supervised vs. unsupervised learning
- Training and test dataset
- Training loss
- Model

# Guidelines to implement Deep Learning schemes



1. Problem formulation (inputs/outputs)
2. Data collection (cf. supervised vs. non-supervised)
3. Definition of performance metrics
4. Selection of neural architectures (at least 2 models)
5. Selection of a training loss
6. Split dataset into training / validation / test datasets
7. Train the selected models from the training dataset and save the best models onto the validation dataset
8. Benchmark the performance of the trained models onto the test dataset
9. Update/iterate 4-5-6-7-8



# Image classification case-study

Let's go

[https://github.com/CIA-Oceanix/DLCourse\\_MOi\\_2022/blob/main/notebooks/notebook\\_MNIST\\_classification\\_MLP\\_with\\_correction.ipynb](https://github.com/CIA-Oceanix/DLCourse_MOi_2022/blob/main/notebooks/notebook_MNIST_classification_MLP_with_correction.ipynb)

# Image classification case-study with pytorch

## 1. Problem formulation (inputs/outputs)




# Image classification case-study with pytorch

## 2. Data collection

```
train_data = datasets.MNIST(root = 'data', train = True, download = True, transform = transform)
test_data = datasets.MNIST(root = 'data', train = False, download = True, transform = transform)
```

## 3. Performance metrics



# Image classification case-study with pytorch

## 4. Neural architecture

```
import torch.nn as nn
import torch.nn.functional as F

class MLP(nn.Module):
    def __init__(self): # FUNCTION TO BE COMPLETED
        super(MLP, self).__init__()
        hidden_1, hidden_2 = 512, 256
        self.fc1 = nn.Linear(28*28, hidden_1)
        self.fc2 = nn.Linear(hidden_1, hidden_2)
        self.fc3 = nn.Linear(hidden_2, 10)
        self.dropout = nn.Dropout(0.2)

    def forward(self, x): # FUNCTION TO BE COMPLETED
        x = x.view(-1, 28*28)
        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = F.relu(self.fc2(x))
        x = self.dropout(x)
        x = self.fc3(x)
        return x
```

## 5. Training loss

```
criterion = nn.CrossEntropyLoss() # TO DO
```

Model complexity ?



# Image classification case-study with pytorch

## 6. Split dataset into training / validation / test datasets

```
batch_size = 20
valid_size = 0.2
train_size = 0.3

def create_data_loaders(batch_size, valid_size, train_data, test_data): # FUNCTION TO BE COMPLETED

    num_train = len(train_data)
    indices = list(range(num_train))
    np.random.shuffle(indices)
    nb_train = int( np.floor(train_size * num_train ) )
    split = int(np.floor(valid_size * num_train))
    train_index, valid_index = indices[split:nb_train], indices[:split]

    train_sampler = SubsetRandomSampler(train_index)
    valid_sampler = SubsetRandomSampler(valid_index)

    train_loader = torch.utils.data.DataLoader(train_data, batch_size = batch_size, sampler = train_sampler)
    valid_loader = torch.utils.data.DataLoader(train_data, batch_size = batch_size, sampler = valid_sampler)
    test_loader = torch.utils.data.DataLoader(test_data, batch_size = batch_size)

    return train_loader, valid_loader, test_loader
```

# Image classification case-study with pytorch

## 7. Model training

```
optimizer = torch.optim.SGD(model_1.parameters(), lr = 0.01)
```

```
n_epochs = 30
```

```
def training(n_epochs, train_loader, valid_loader, model, criterion, optimizer):
```

```
    train_losses, valid_losses = [], []
```

```
    valid_loss_min = np.Inf
```

```
    for epoch in range(n_epochs):
```

```
        train_loss, valid_loss = 0, 0
```

```
        model.train()
```

```
        for data, label in train_loader:
```

```
            data = data.to(device=device, dtype=torch.float32)
```

```
            label = label.to(device=device, dtype=torch.long)
```

```
            optimizer.zero_grad()
```

```
            output = model(data)
```

```
            loss = criterion(output, label)
```

```
            loss.backward()
```

```
            optimizer.step()
```

```
            train_loss += loss.item() * data.size(0)
```

# Image classification case-study

Go and run the notebook

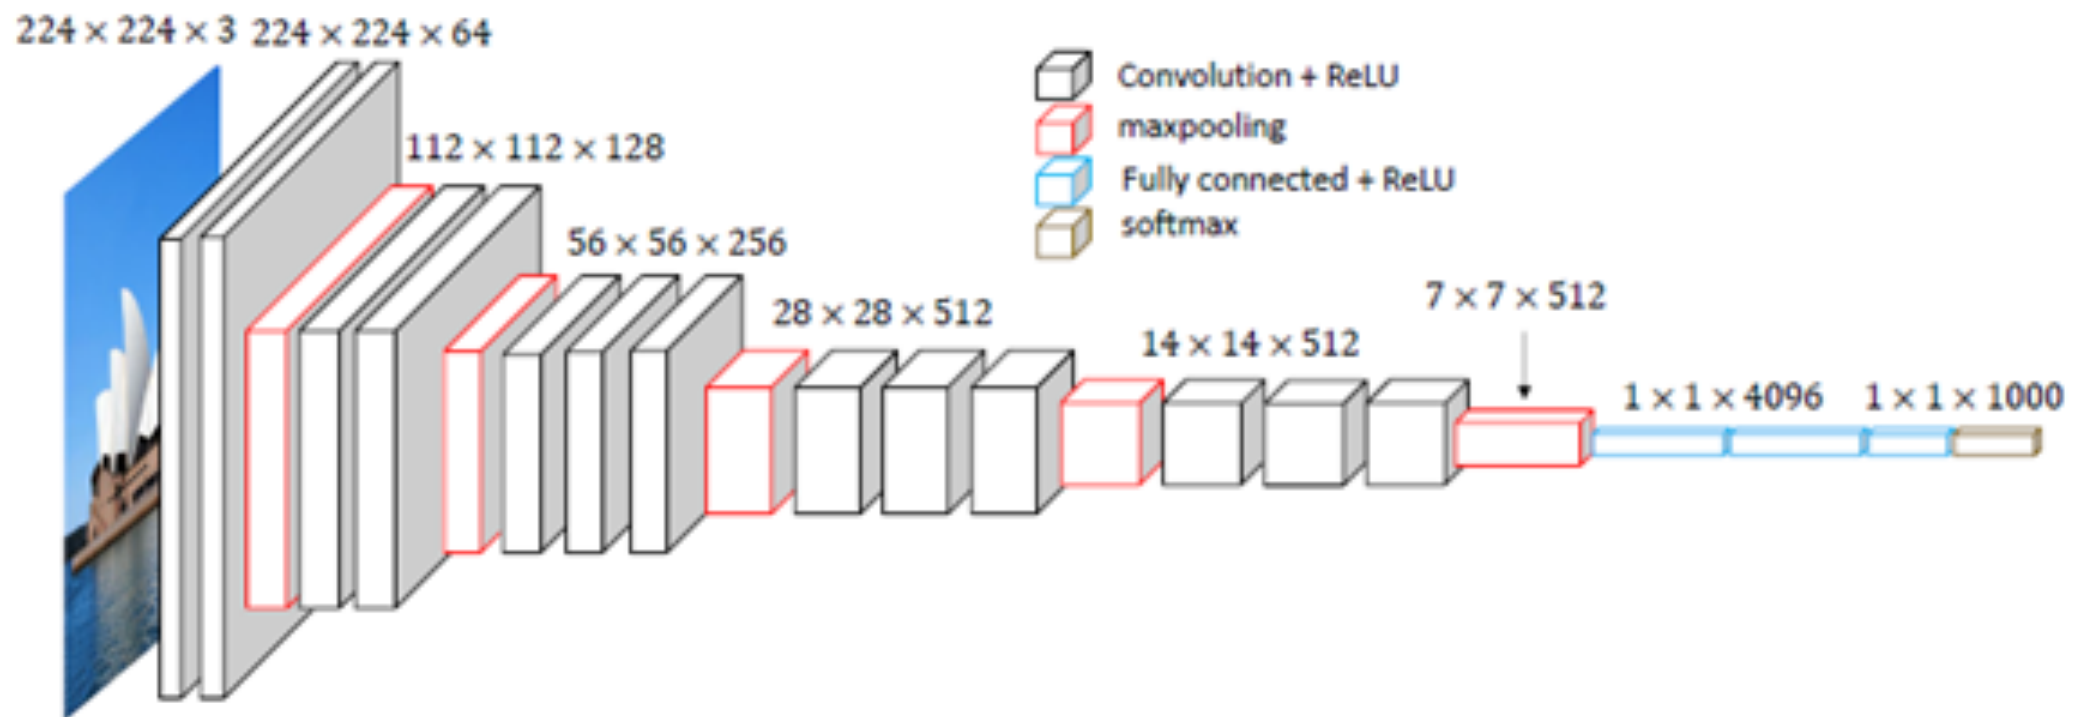
Questions :

What is the effect of the dropout layer in the MLP architecture ?

# Convolutional Neural Networks

# State-of-the-art NNs in computer vision

DL models are (in general) feedforward models. VGG16 as an illustration



**Elementary components**

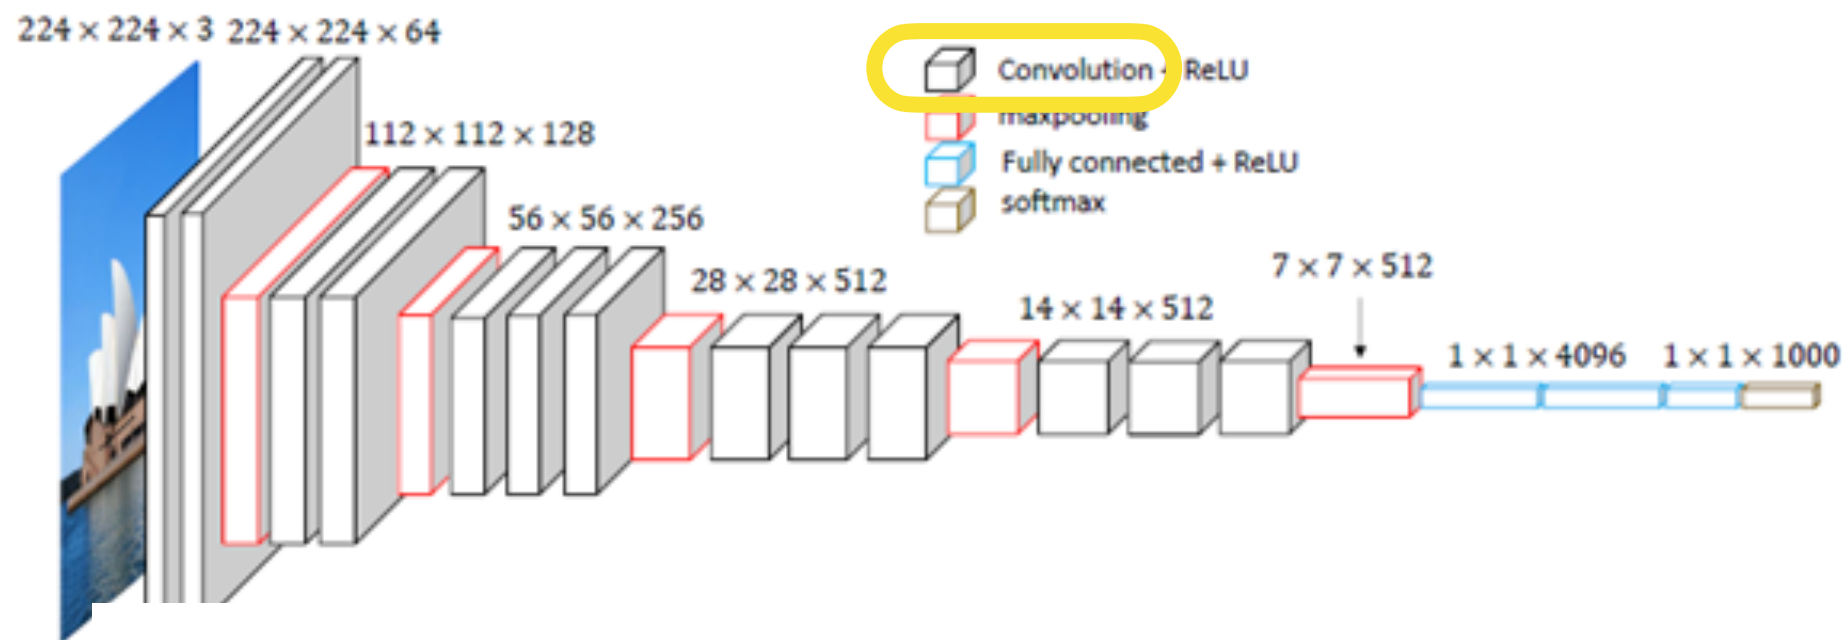
Convolution layers

Activation layers

Pooling layers

Dense layers

# Basics of DL models




Elementary  
components

Convolution layers

Activation layers

Pooling layers

Dense layers




```
torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```


<https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

Number of parameters ?  
Independent on the sizes of the input  
and output layer

# Dense layer vs Conv layer



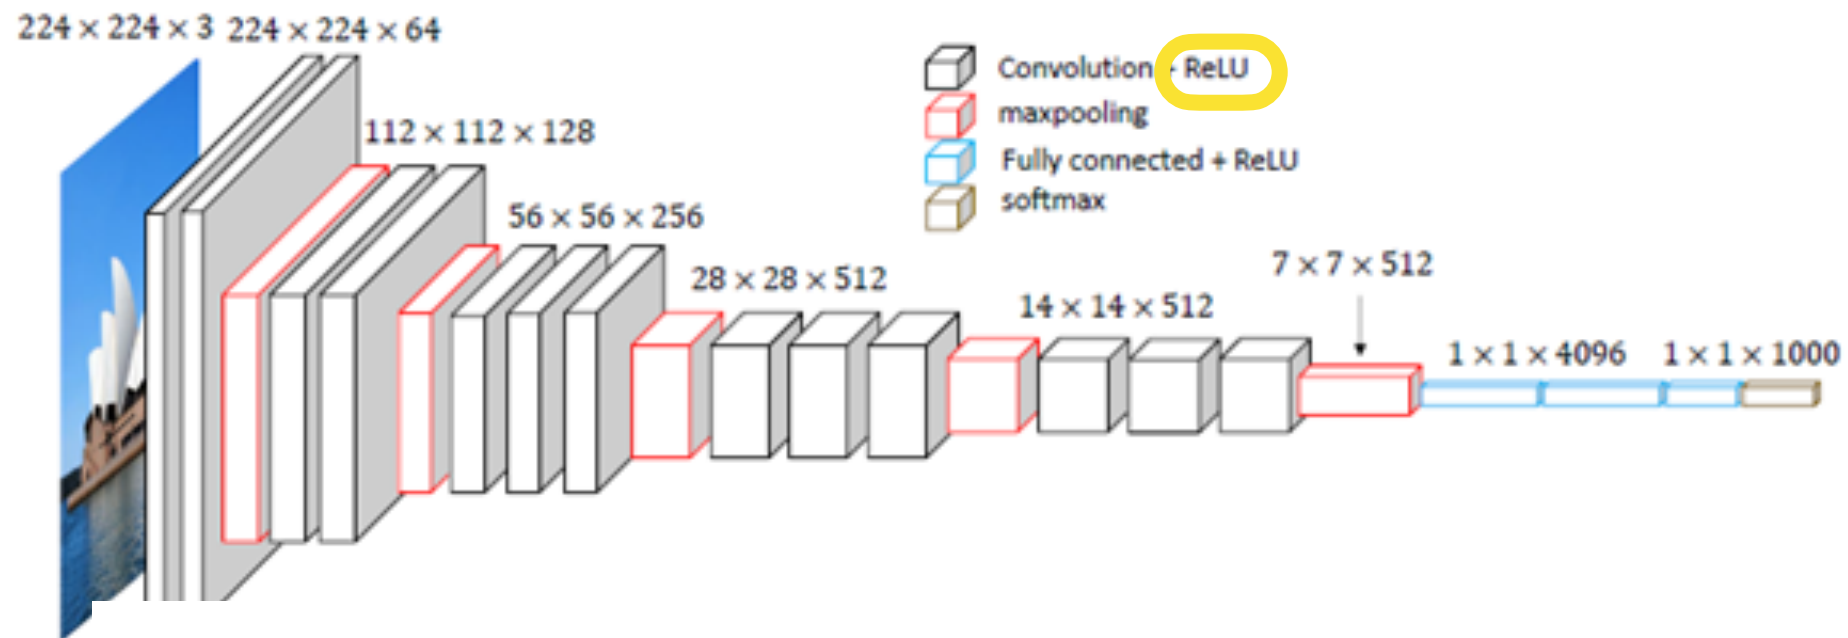
**Dense layer**



**Convolutional layer**



# Basics of DL models



Elementary components



Convolution layers

Activation layers

Pooling layers


Dense layers

ReLU  
(Rectified  
Linear Unit)





# Basics of DL models



Elementary components


Convolution layers

Activation layers


Pooling layers

Dense layers

An example of max-pooling operator




Pooling downsamples the input layer



```
torch.nn.MaxPool2d(kernel_size, stride=None, padding=0, dilation=1, return_indices=False, ceil_mode=False) [SOURCE]
```

# Basics of DL models



Elementary  
components


Convolution layers

Activation layers



Pooling layers

Dense layers

Dense layers  
or  
Fully-connected (FC) layer  
as in a classic MLP

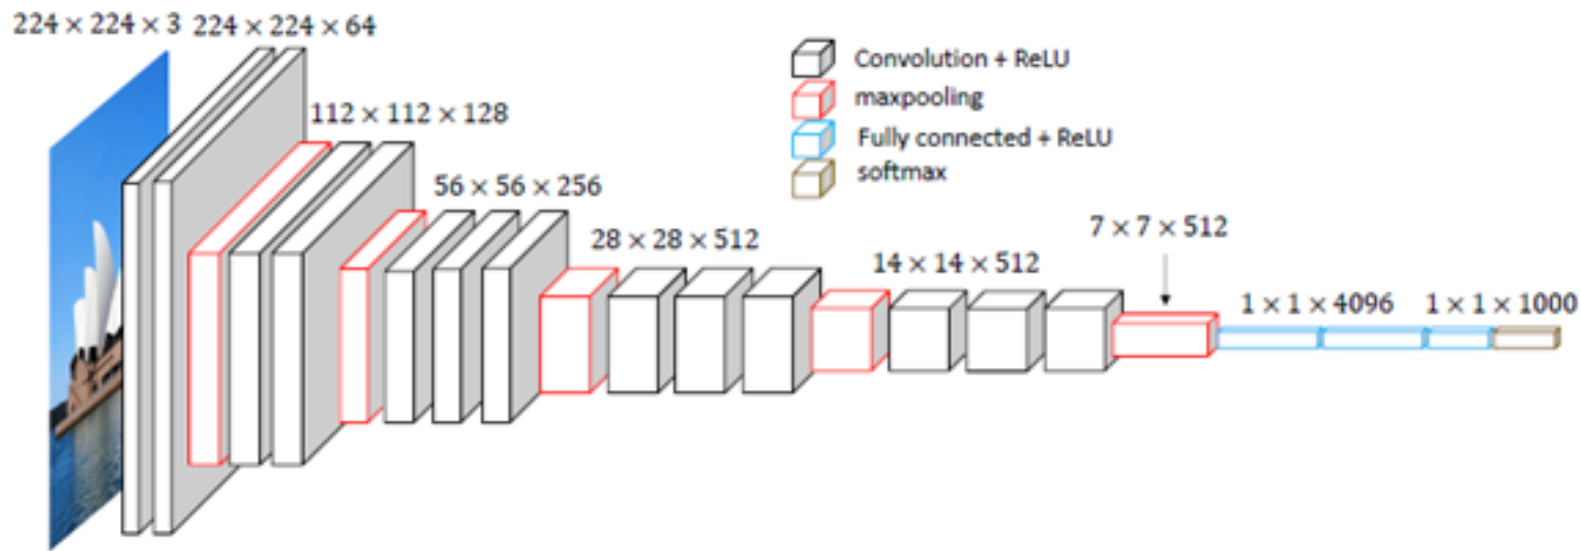


# Image classification case-study with CNN




[https://github.com/CIA-Oceanix/DLCourse\\_MOi\\_2022/blob/main/notebooks/notebook\\_MNIST\\_classification\\_MLP\\_CNN\\_students.ipynb](https://github.com/CIA-Oceanix/DLCourse_MOi_2022/blob/main/notebooks/notebook_MNIST_classification_MLP_CNN_students.ipynb)

# Examples of DL models for object recognition (2010-2020)





## VGG16

( $<100\text{M}$  of parameters)




**(5M of parameters)**



# AlexNet

(60M of parameters)

# DL and Benchmarking (Data Challenges)



<https://paperswithcode.com/sota/image-classification-on-imagenet>



# of object classes: 1000

# of images > 1.2 M

Best accuracy score: ~91%

State-of-the-art architectures: CNN, Vision Transformers



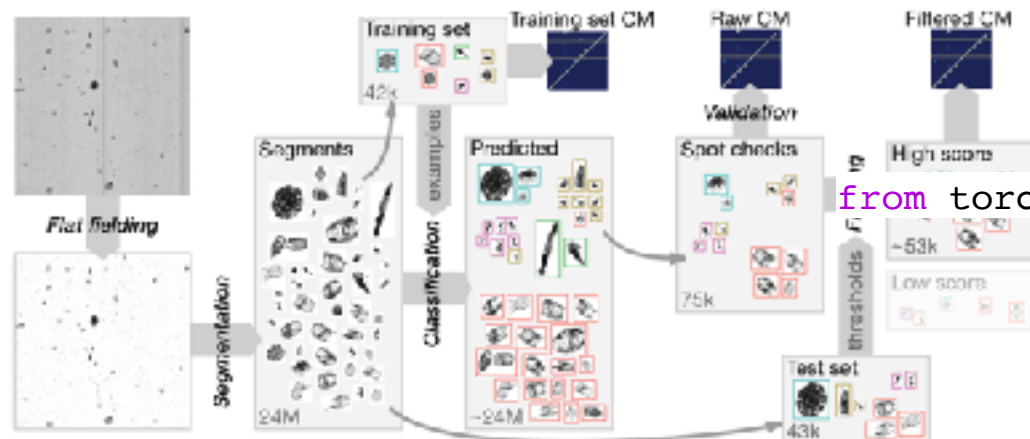
# CNN-based classification and Ocean Data

LIMNOLOGY  
and  
OCEANOGRAPHY: METHODS

## Automated plankton image analysis using convolutional neural networks



Jessica Y. Luo<sup>1,2,\*</sup>, Jean-Olivier Irisson<sup>3</sup>, Benjamin Graham<sup>4</sup>, Cedric Guigand<sup>1</sup>, Amin Sarafraz<sup>5</sup>

Christi  
<sup>1</sup>Marine  
<sup>2</sup>Harfvel  
<sup>3</sup>Soeben  
<sup>4</sup>Depart  
<sup>5</sup>Center



## Tropical Cyclone Intensity Classification and Estimation Using Infrared Satellite Images With Deep Learning

Chang-Jiang Zhang<sup>1,\*</sup>, Xian-Ji




ASLO

Journal of the  
© 2018 Association for the Science of Limnology and Oceanography  
doi: 10.1002/lom.10100

IEEE TRANSACTIONS ON GEOSCIENCE AND REMOTE SENSING

## Vertical Structure-Based Classification of Oceanic Eddy Using 3-D Convolutional Neural Network

Baoxiang Huang<sup>1,\*</sup>



Contents lists available at ScienceDirect

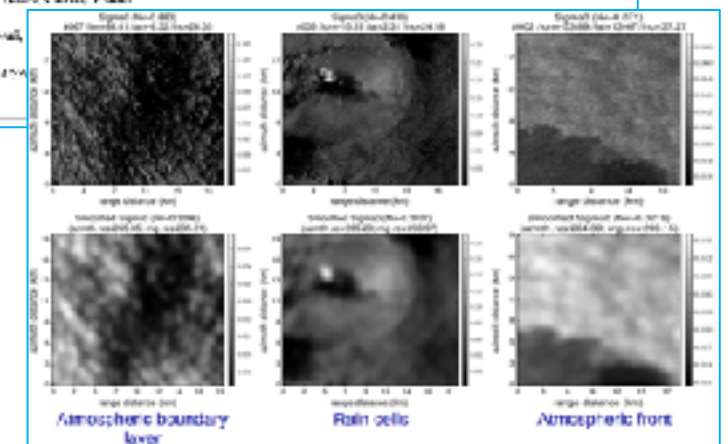
Remote Sensing of Environment

journal homepage: [www.elsevier.com/locate/rse](http://www.elsevier.com/locate/rse)

Classification of the global Sentinel-1 SAR vignettes for ocean surface process studies

Chen Wang<sup>a,b,\*</sup>, Pierre Tandoi<sup>b</sup>, Alexis Mouche<sup>a</sup>, Justin E. Stopa<sup>c</sup>, Victor Gressani<sup>a</sup>, Nicolas Longepe<sup>c</sup>, Douglas Vandemark<sup>c</sup>, Ralph C. Proster<sup>d</sup>, Bertrand Chapron<sup>a</sup>

<sup>a</sup>IFREMER, Univ. Brest, CNRS, IRD, Laboratoire d'Océanographie Physique et Spatial (LOGOS), Brest, France  
<sup>b</sup>IMT Atlantique, Lab-STICC, UMR, Brest, France  
<sup>c</sup>Department of Ocean Sciences and Engineering, University of Miami at Miami, Miami, FL  
<sup>d</sup>Space and Ground Segment, College de l'Atlantique Sud (CLAS), Rimouski, France  
\*Corresponding author. E-mail address: chen.wang@ifremer.fr (C. Wang).  
Applied Marine Technology, University of Technology, Sydney, NSW, Australia



# Fine-tuning from pre-trained models

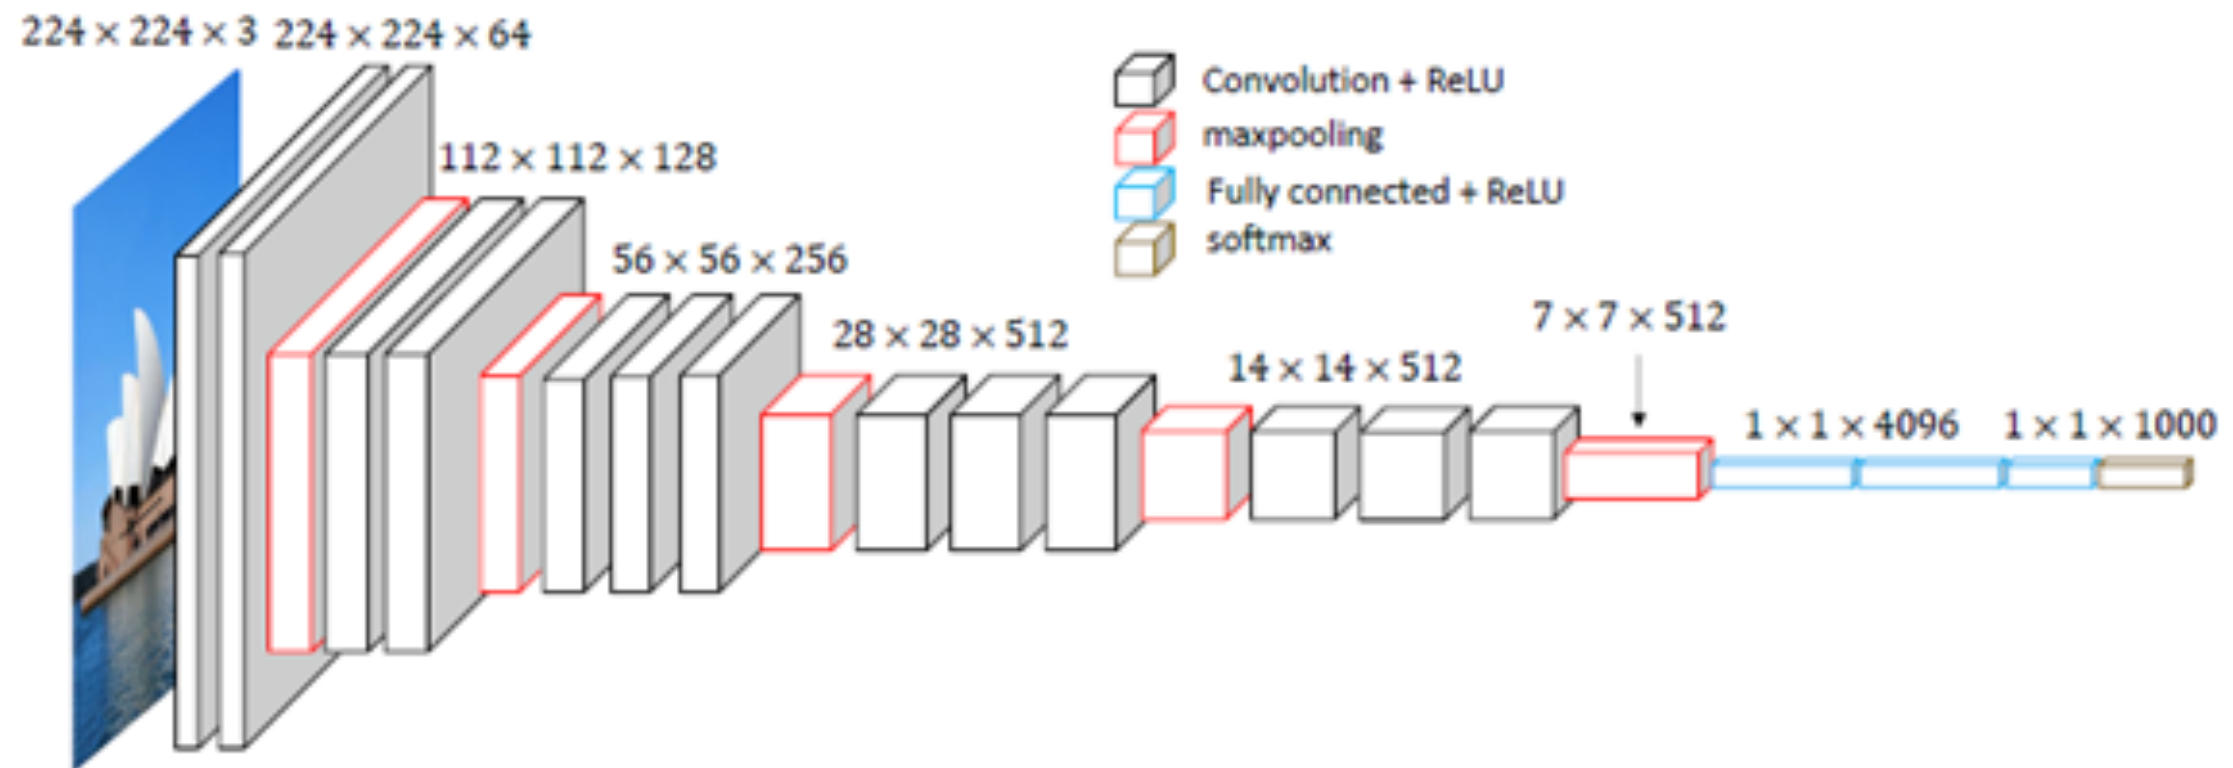
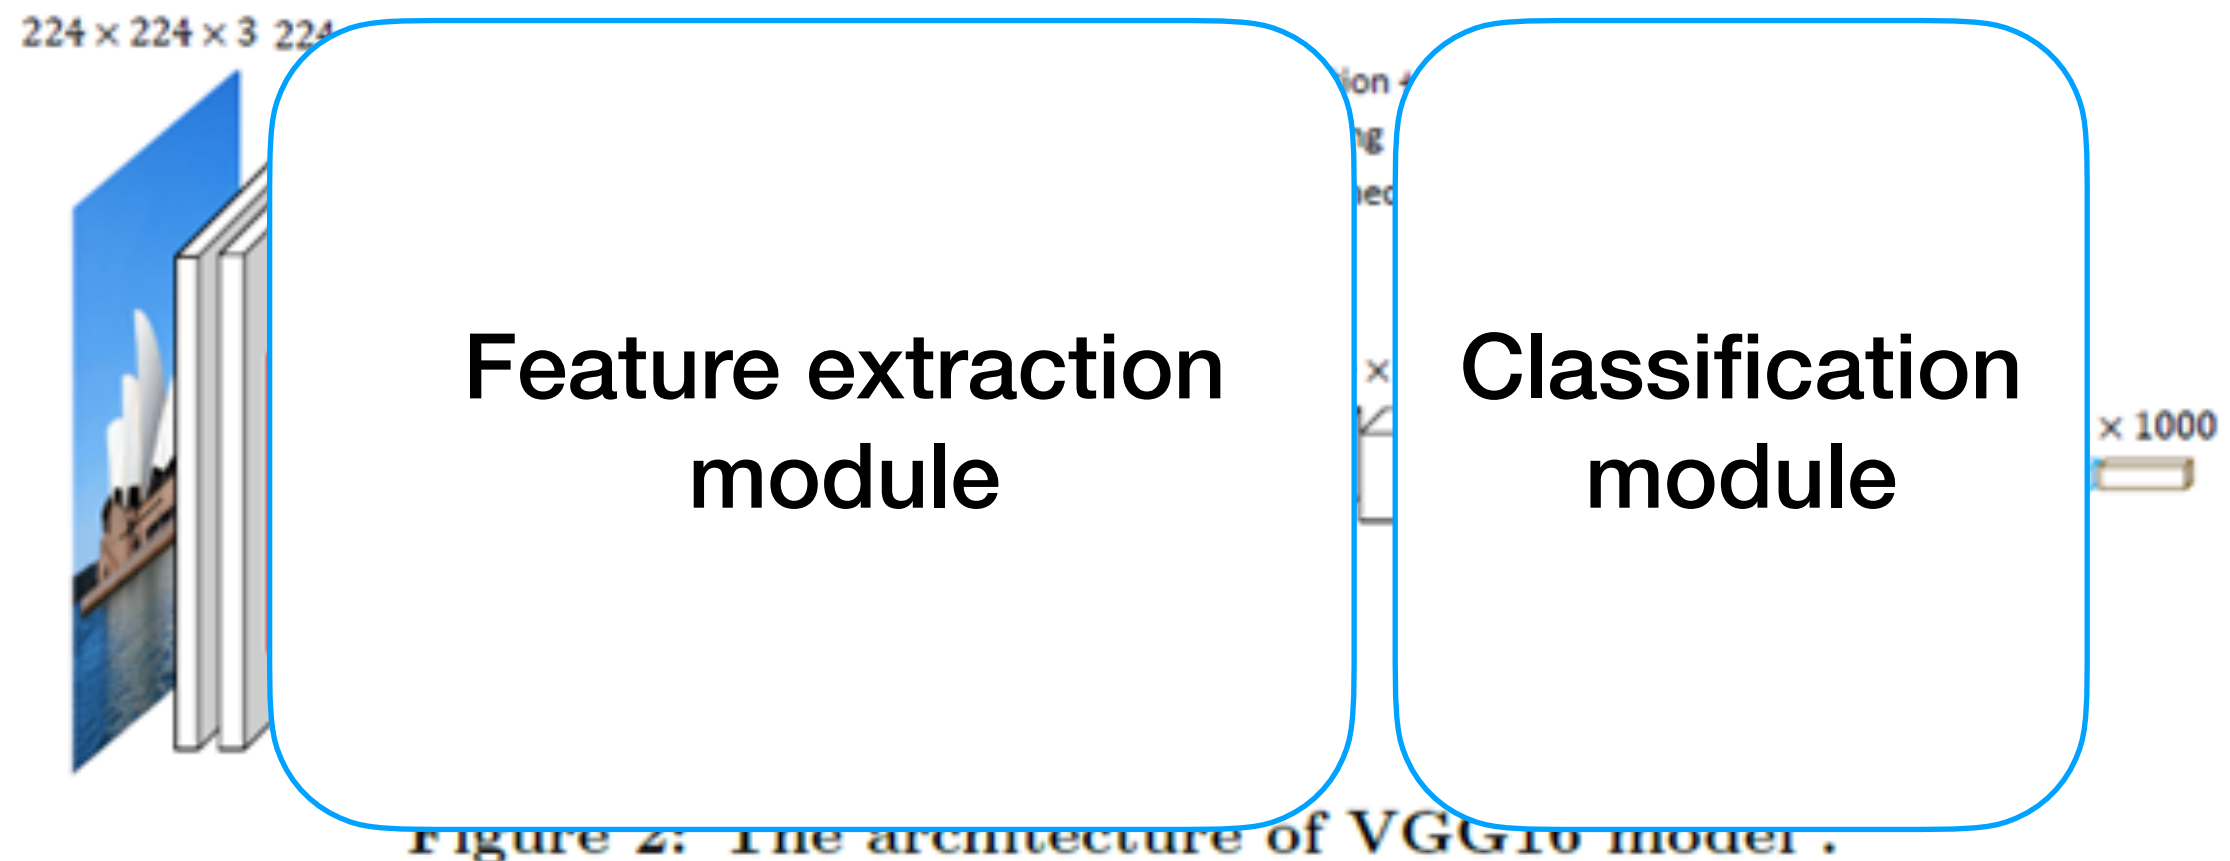


Figure 2: The architecture of VGG16 model .

**General idea:** the first layers involve generic feature extraction step and the last block can be regarded as a dataset-specific classification block.

# Fine-tuning from pre-trained models



**General idea:** the first layers involve generic feature extraction step and the last block can be regarded as a dataset-specific classification block.



# Fine-tuning from pre-trained models

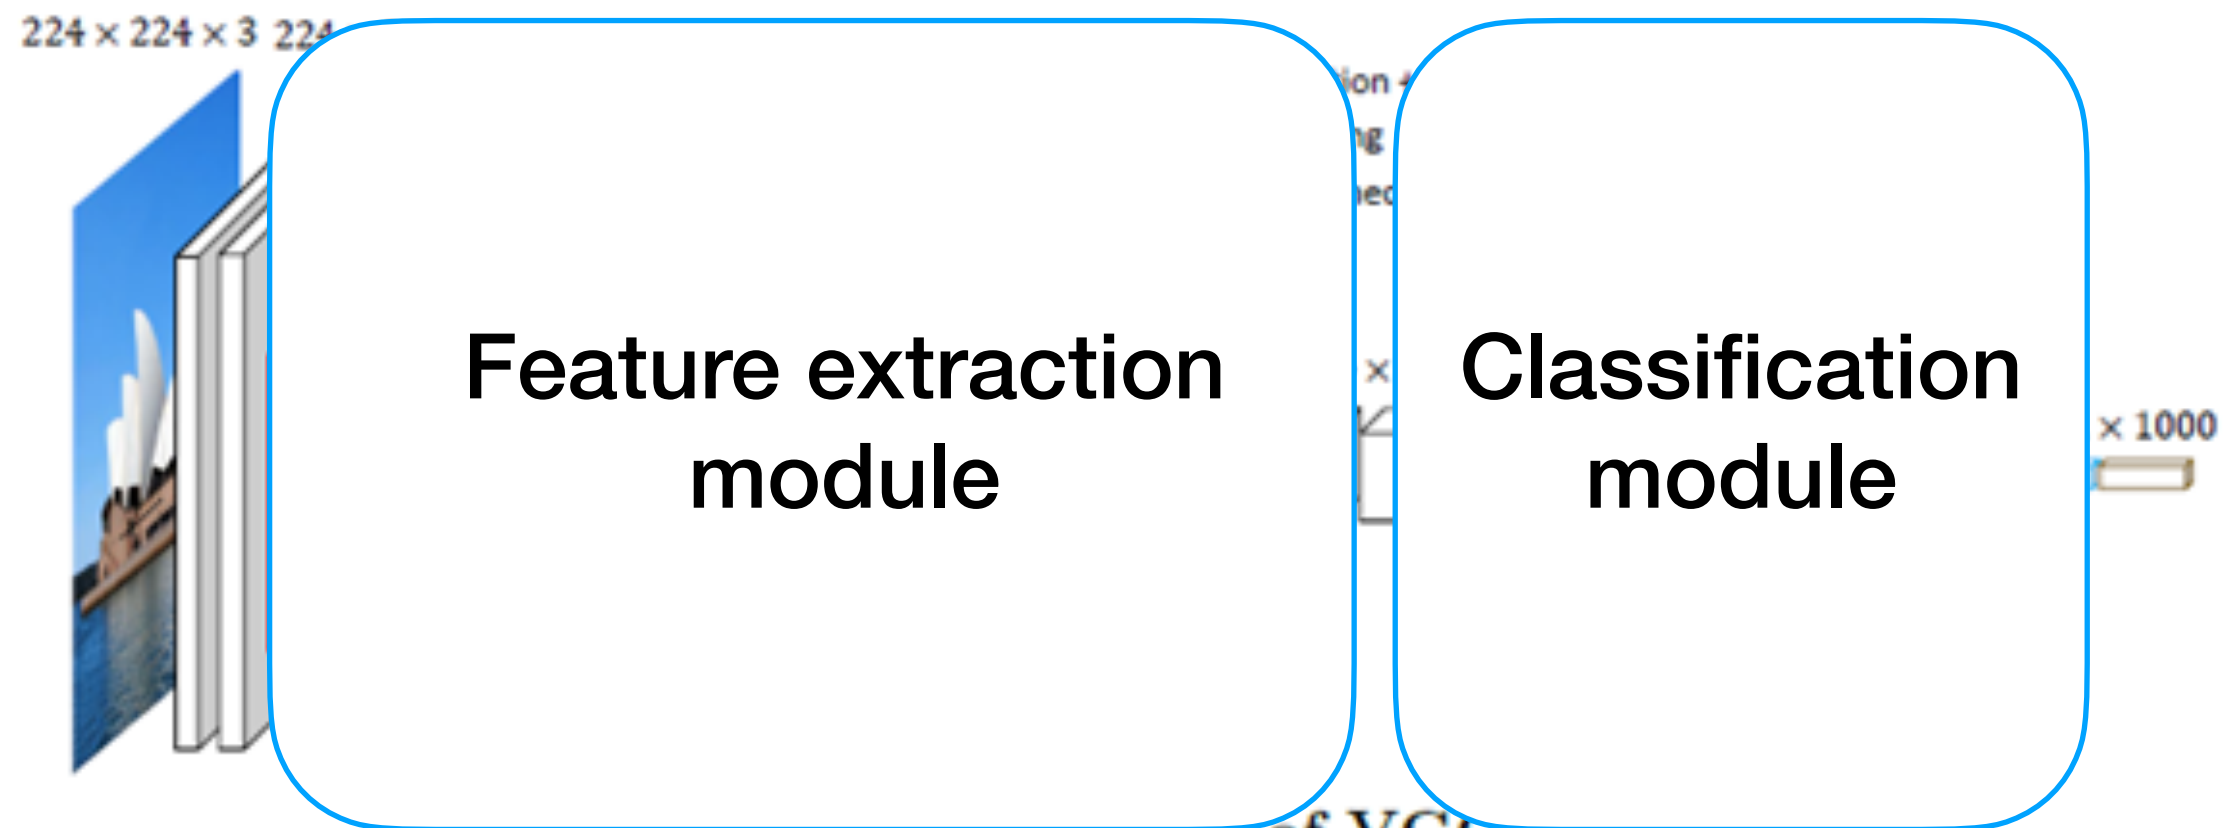




Figure 4: The architecture of VGG16 model.

[https://github.com/CIA-Oceanix/DLCourse\\_MOi\\_2022/blob/main/notebooks/notebook\\_MNIST\\_classification\\_MLP\\_CNN\\_TransferLearning\\_students.ipynb](https://github.com/CIA-Oceanix/DLCourse_MOi_2022/blob/main/notebooks/notebook_MNIST_classification_MLP_CNN_TransferLearning_students.ipynb)

# Over-fitting



# Over-fitting



# Regularization tricks to avoid overfitting

- Penalty terms In the training loss
- Data augmentation
- Dropout layers

# Parameter norm penalization

- Regularized objective function:

$$\tilde{J}(\theta) = J(\theta) + \alpha\Omega(\theta)$$

- L<sup>2</sup> norm:  $\Omega(\theta) = \frac{1}{2}||w||_2^2$
- L<sup>1</sup> norm:  $\Omega(\theta) = ||w||_1 = \sum_i |w_i|$

# Data augmentation

- Purpose: improving model generalization error by training on more data
- Very efficient for object recognition
- How to:
  - apply (geometric) transformations on input data (such as translation, rotation, scaling for images).
  - noise injection

# Dropout





Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

# A rapid tour of the CNN zoo

- CNN
- Convolution auto-encoders
  - U-Nets
  - Residual Networks




# Examples of DL models for object recognition (2010-2020)



**VGG16**  
(<100M of parameters)

**Google Inception**

(5M of parameters)



**AlexNet**  
(60M of parameters)


# Lecture. #2

## Things to know

- Convolution layers
- Pooling layers
- Activation layers
- Dropout layers
- Padding and stride
- Fine-tuning
- Over-fitting
- Data augmentation
-

# Auto-encoders


# Auto-encoders




Output with the same shape as the input

Application ?

# Dense auto-encoders




# PCA/EOF



PCA as a linear auto-encoder architecture.


Which additional constraint ?

# Convolutional auto-encoders




# Convolutional AE Zoo

Many applications do not require a low-dimensional representation  
(e.g., denoising, interpolation, super-resolution,...)



<https://arxiv.org/pdf/1606.08921.pdf>




UNet

Ronneberger  
et al., 2015



# Convolutional AE Zoo



Residual Block  
(ResNet)

Rousseau et al.,  
2019

Often used to address vanishing gradients (“very” deep networks)