

Course #2:

**Deep Learning, from
MLP to CNN**

Roadmap


- Recap from course #1
- MLP and Image classification as a case study
- CNN: basic principles
- Application to image classification
- Classic CNN architectures
- Auto-encoders

Recap from Course #1

Things to know

- Supervised vs. unsupervised learning
- Training and test dataset
- Training loss
- Model

Guidelines to implement Deep Learning schemes



1. Problem formulation (inputs/outputs)
2. Data collection (cf. supervised vs. non-supervised)
3. Definition of performance metrics
4. Selection of neural architectures (at least 2 models)
5. Selection of a training loss
6. Split dataset into training / validation / test datasets
7. Train the selected models from the training dataset and save the best models onto the validation dataset
8. Benchmark the performance of the trained models onto the test dataset
9. Update/iterate 4-5-6-7-8

Image classification case-study

Let's go

https://github.com/CIA-Oceanix/DLCourse_MOi_2022/blob/main/notebooks/notebook_MNIST_classification_MLP_with_correction.ipynb

Image classification case-study with pytorch

1. Problem formulation (inputs/outputs)





Image classification case-study with pytorch

2. Data collection

```
train_data = datasets.MNIST(root = 'data', train = True, download = True, transform = transform)
test_data = datasets.MNIST(root = 'data', train = False, download = True, transform = transform)
```

3. Performance metrics




Image classification case-study with pytorch

4. Neural architecture

```
import torch.nn as nn
import torch.nn.functional as F

class MLP(nn.Module):
    def __init__(self): # FUNCTION TO BE COMPLETED
        super(MLP, self).__init__()
        hidden_1, hidden_2 = 512, 256
        self.fc1 = nn.Linear(28*28, hidden_1)
        self.fc2 = nn.Linear(hidden_1, hidden_2)
        self.fc3 = nn.Linear(hidden_2, 10)
        self.dropout = nn.Dropout(0.2)

    def forward(self, x): # FUNCTION TO BE COMPLETED
        x = x.view(-1, 28*28)
        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = F.relu(self.fc2(x))
        x = self.dropout(x)
        x = self.fc3(x)
        return x
```

5. Training loss

```
criterion = nn.CrossEntropyLoss() # TO DO
```

Model complexity ?

Image classification case-study with pytorch

6. Split dataset into training / validation / test datasets

```
batch_size = 20
valid_size = 0.2
train_size = 0.3

def create_data_loaders(batch_size, valid_size, train_data, test_data): # FUNCTION TO BE COMPLETED

    num_train = len(train_data)
    indices = list(range(num_train))
    np.random.shuffle(indices)
    nb_train = int( np.floor(train_size * num_train ))
    split = int(np.floor(valid_size * num_train))
    train_index, valid_index = indices[split:nb_train], indices[:split]

    train_sampler = SubsetRandomSampler(train_index)
    valid_sampler = SubsetRandomSampler(valid_index)

    train_loader = torch.utils.data.DataLoader(train_data, batch_size = batch_size, sampler = train_sampler)
    valid_loader = torch.utils.data.DataLoader(train_data, batch_size = batch_size, sampler = valid_sampler)
    test_loader = torch.utils.data.DataLoader(test_data, batch_size = batch_size)

    return train_loader, valid_loader, test_loader
```

Image classification case-study with pytorch

7. Model training

```
optimizer = torch.optim.SGD(model_1.parameters(), lr = 0.01)
```

```
n_epochs = 30
```

```
def training(n_epochs, train_loader, valid_loader, model, criterion, optimizer):
```

```
    train_losses, valid_losses = [], []
```

```
    valid_loss_min = np.Inf
```

```
    for epoch in range(n_epochs):
```

```
        train_loss, valid_loss = 0, 0
```

```
        model.train()
```

```
        for data, label in train_loader:
```

```
            data = data.to(device=device, dtype=torch.float32)
```

```
            label = label.to(device=device, dtype=torch.long)
```

```
            optimizer.zero_grad()
```

```
            output = model(data)
```

```
            loss = criterion(output, label)
```

```
            loss.backward()
```

```
            optimizer.step()
```

```
            train_loss += loss.item() * data.size(0)
```

Image classification case-study

Go and run the notebook

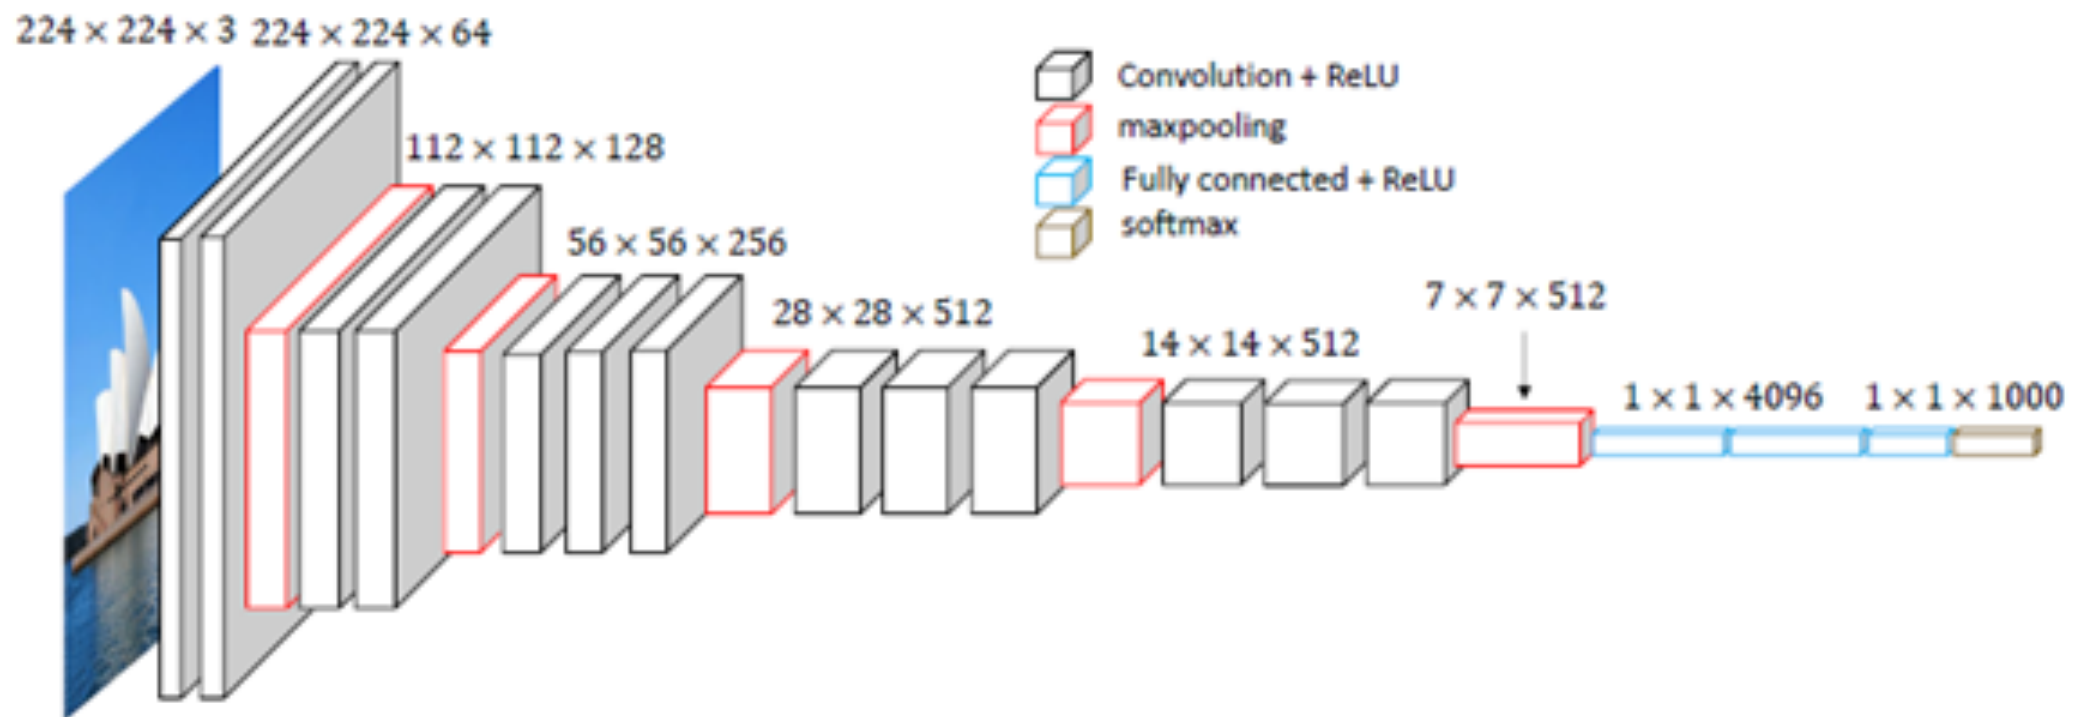
Questions :

What is the effect of the dropout layer in the MLP architecture ?

Convolutional Neural Networks

State-of-the-art NNs in computer vision

DL models are (in general) feedforward models. VGG16 as an illustration



Elementary components

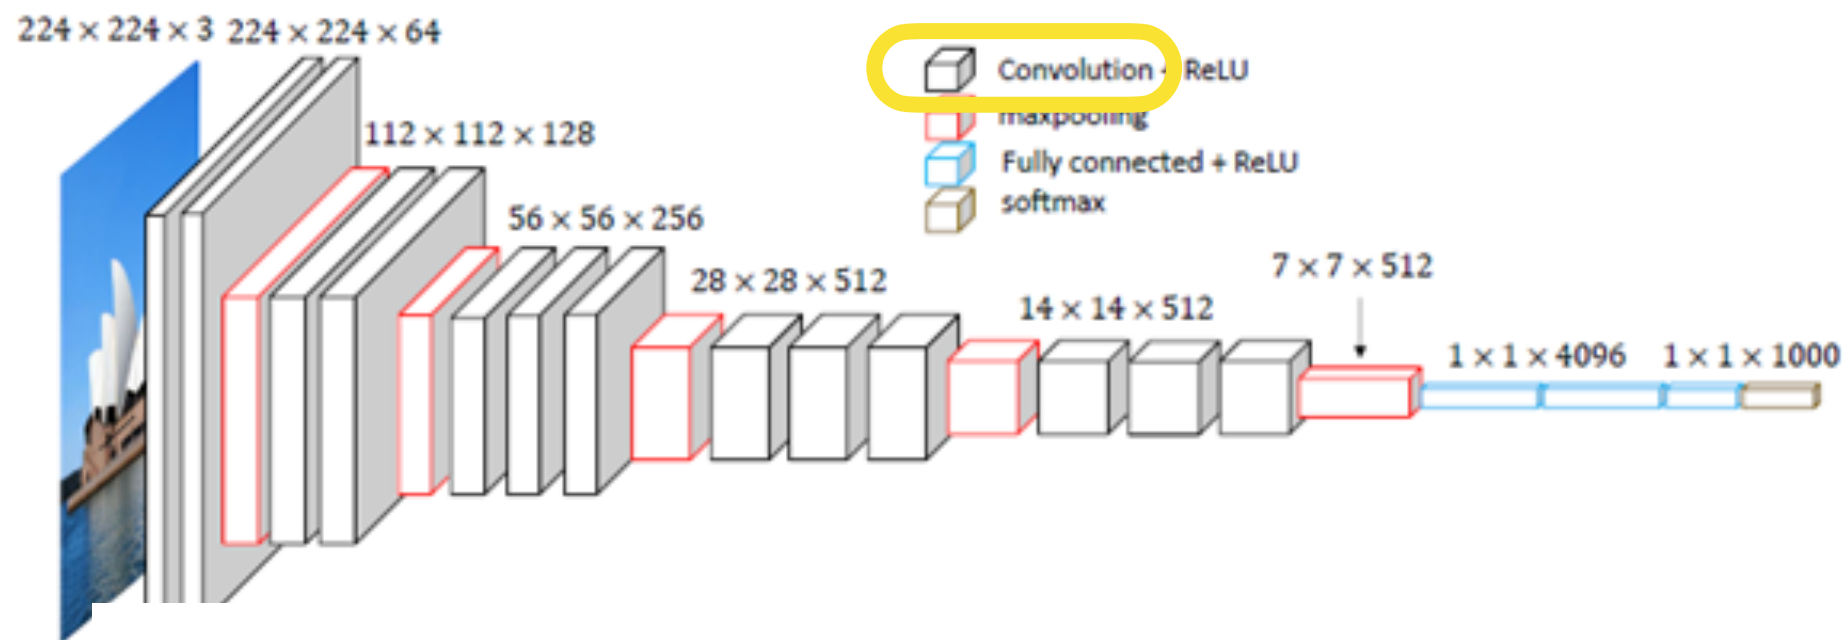
Convolution layers

Activation layers

Pooling layers

Dense layers

Basics of DL models




Elementary
components

Convolution layers

Activation layers

Pooling layers

Dense layers




```
torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```


<https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

Number of parameters ?
Independent on the sizes of the input
and output layer

Dense layer vs Conv layer

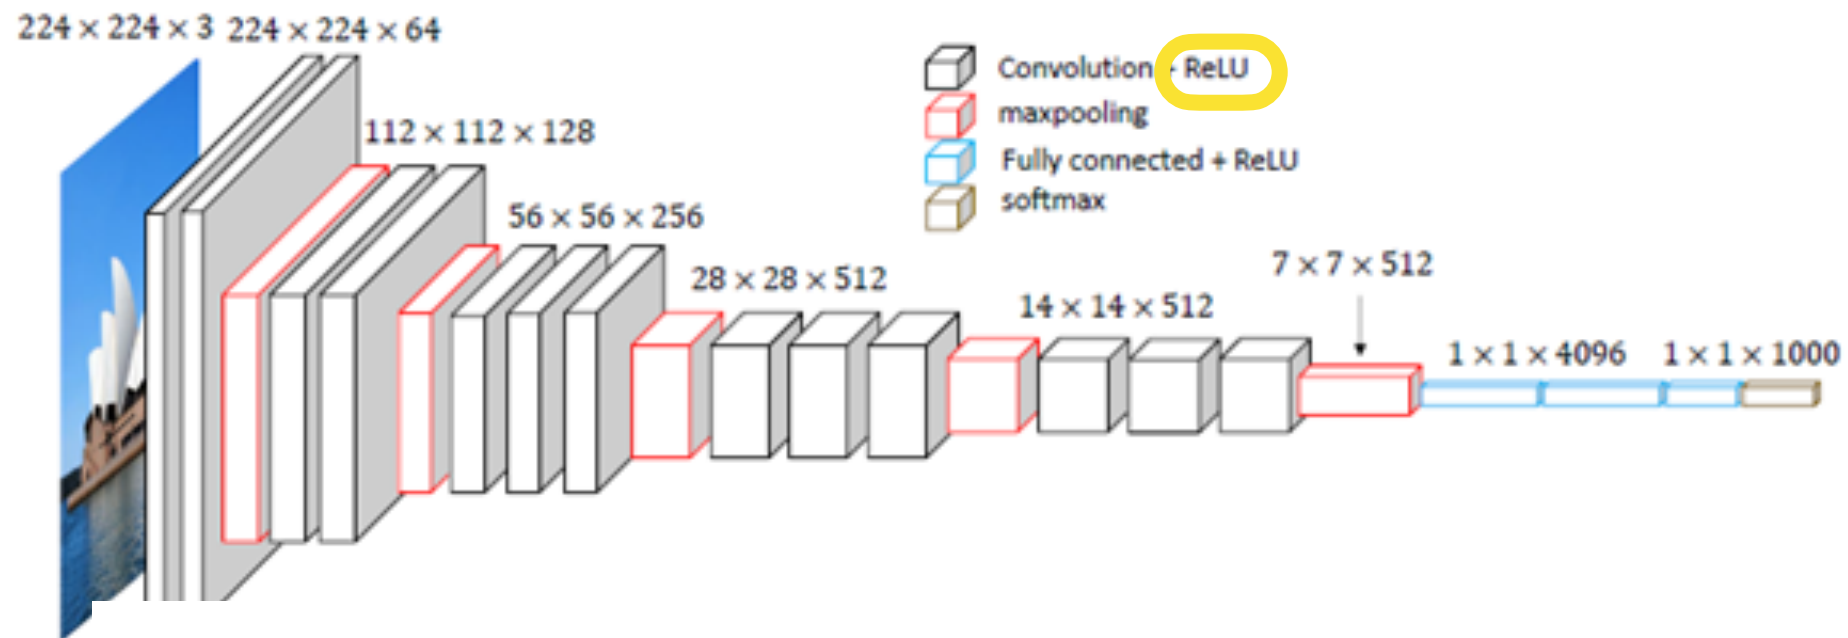


Dense layer



Convolutional layer

Basics of DL models



Elementary
components



Convolution layers

Activation layers


Pooling layers

Dense layers

ReLU
(Rectified
Linear Unit)



Basics of DL models



Elementary components


Convolution layers

Activation layers


Pooling layers

Dense layers

An example of max-pooling operator




Pooling downsamples the input layer



```
torch.nn.MaxPool2d(kernel_size, stride=None, padding=0, dilation=1, return_indices=False, ceil_mode=False) [SOURCE]
```

Basics of DL models



Elementary
components

Convolution layers

Activation layers

Pooling layers

Dense layers

Dense layers
or
Fully-connected (FC) layer
as in a classic MLP




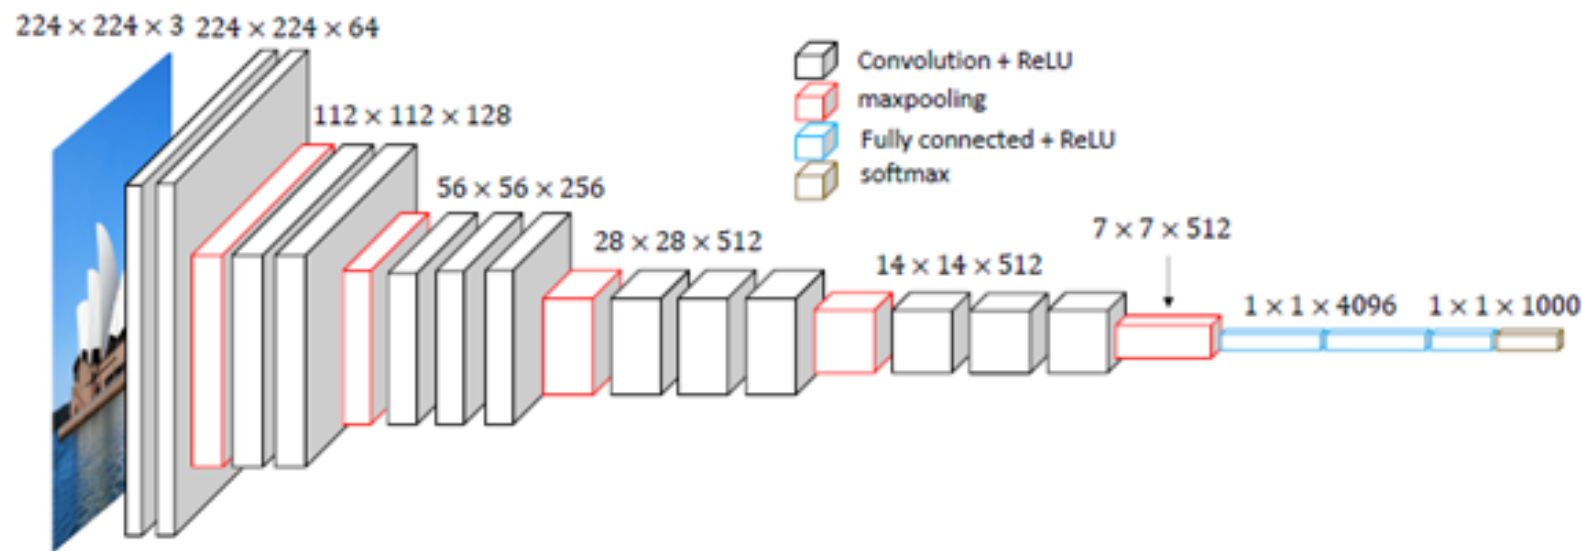


Image classification case-study with CNN



https://github.com/CIA-Oceanix/DLCourse_MOi_2022/blob/main/notebooks/notebook_MNIST_classification_MLP_CNN_students.ipynb


Examples of DL models for object recognition (2010-2020)



VGG16
($<100\text{M}$ of parameters)


Google Inception

(5M of parameters)



AlexNet
(60M of parameters)

DL and Benchmarking (Data Challenges)



<https://paperswithcode.com/sota/image-classification-on-imagenet>



of object classes: 1000

of images > 1.2 M

Best accuracy score: ~91%

State-of-the-art architectures: CNN, Vision Transformers

CNN-based classification and Ocean Data

LIMNOLOGY
and
OCEANOGRAPHY: METHODS

Automated plankton image analysis using convolutional neural networks

Jessica Y. Luo^{1,2,*}, Jean-Olivier Irisson³, Benjamin Graham⁴, Cedric Guigand¹, Amin Sarafraz⁵

Christi

¹Marine

²Harfvel

³Seoben

⁴Depart

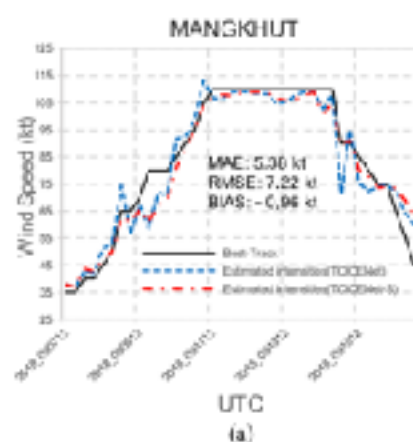
⁵Center



from torchsummary import summary

Tropical Cyclone Intensity Classification and Estimation Using Infrared Satellite Images With Deep Learning

Chang-Jiang Zhang^{1,*}, Xian-Ji




ASLO

Journal of the
© 2018 Association for the Science of Limnology and Oceanography
doi: 10.1002/lom.10100

IEEE TRANSACTIONS ON GEOSCIENCE AND REMOTE SENSING

Vertical Structure-Based Classification of Oceanic Eddy Using 3-D Convolutional Neural Network

Baoxiang Huang^{1,*}



Contents lists available at ScienceDirect

Remote Sensing of Environment

journal homepage: www.elsevier.com/locate/rse

Classification of the global Sentinel-1 SAR vignettes for ocean surface process studies

Chen Wang^{1,*}, Pierre Tandoi², Alexis Mouche³, Justin E. Stopa⁴, Victor Gressani⁵, Nicolas Longepe⁶, Douglas Vandemark⁷, Ralph C. Proster⁸, Bertrand Chapron⁹

¹IFREMER, Univ. Bourg, CMRS, IRD, Laboratoire d'Océanographie Physique et Spatial (LOGPS), Brest, France

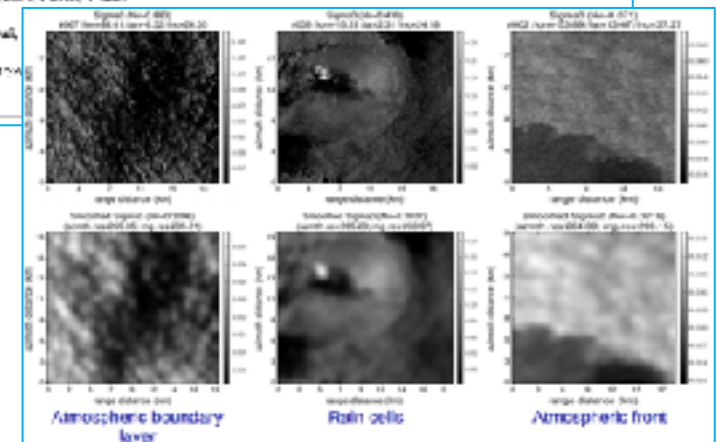
²IMT Atlantique, Lab-STICC, UMR, Brest, France

³Department of Ocean Sciences and Engineering, University of Miami at Miami, Miami, FL

⁴Space and Ground Segment, College de l'Arctique (CSA), Roussay, France

⁵French Research Institute for Exploitation of the Sea (IFREMER), Plouzané, France

⁶Applied Marine Technology, University of Bordeaux, Bordeaux, France



Fine-tuning from pre-trained models

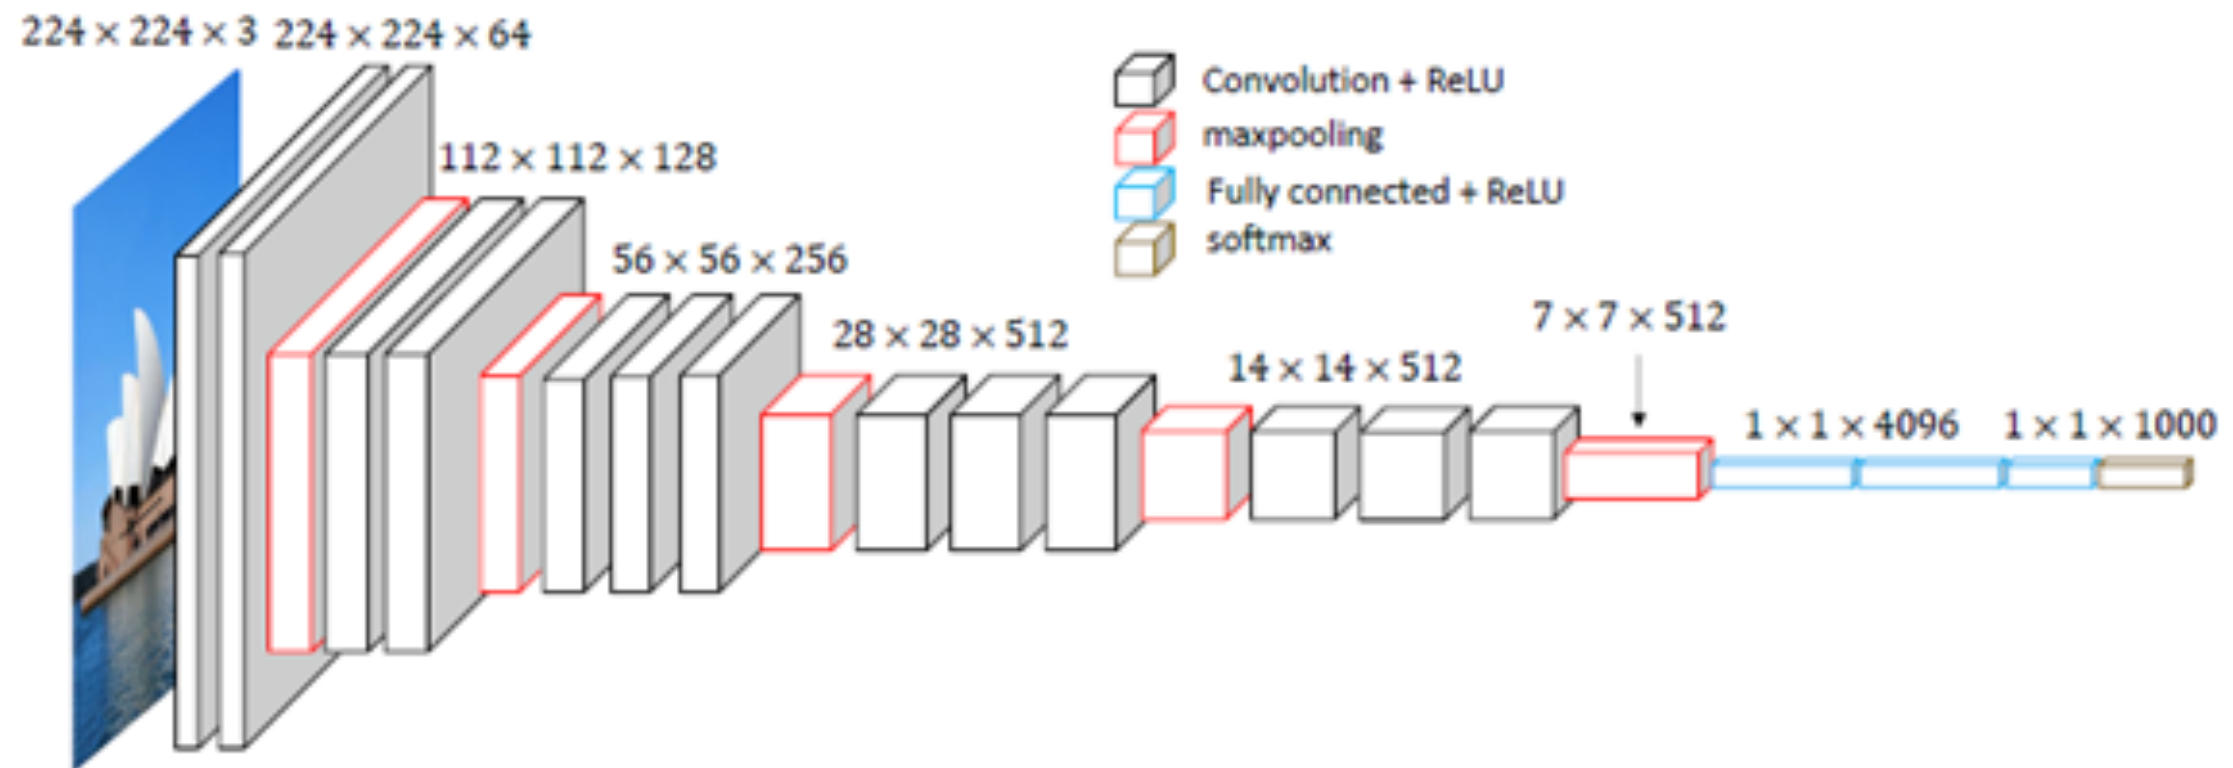
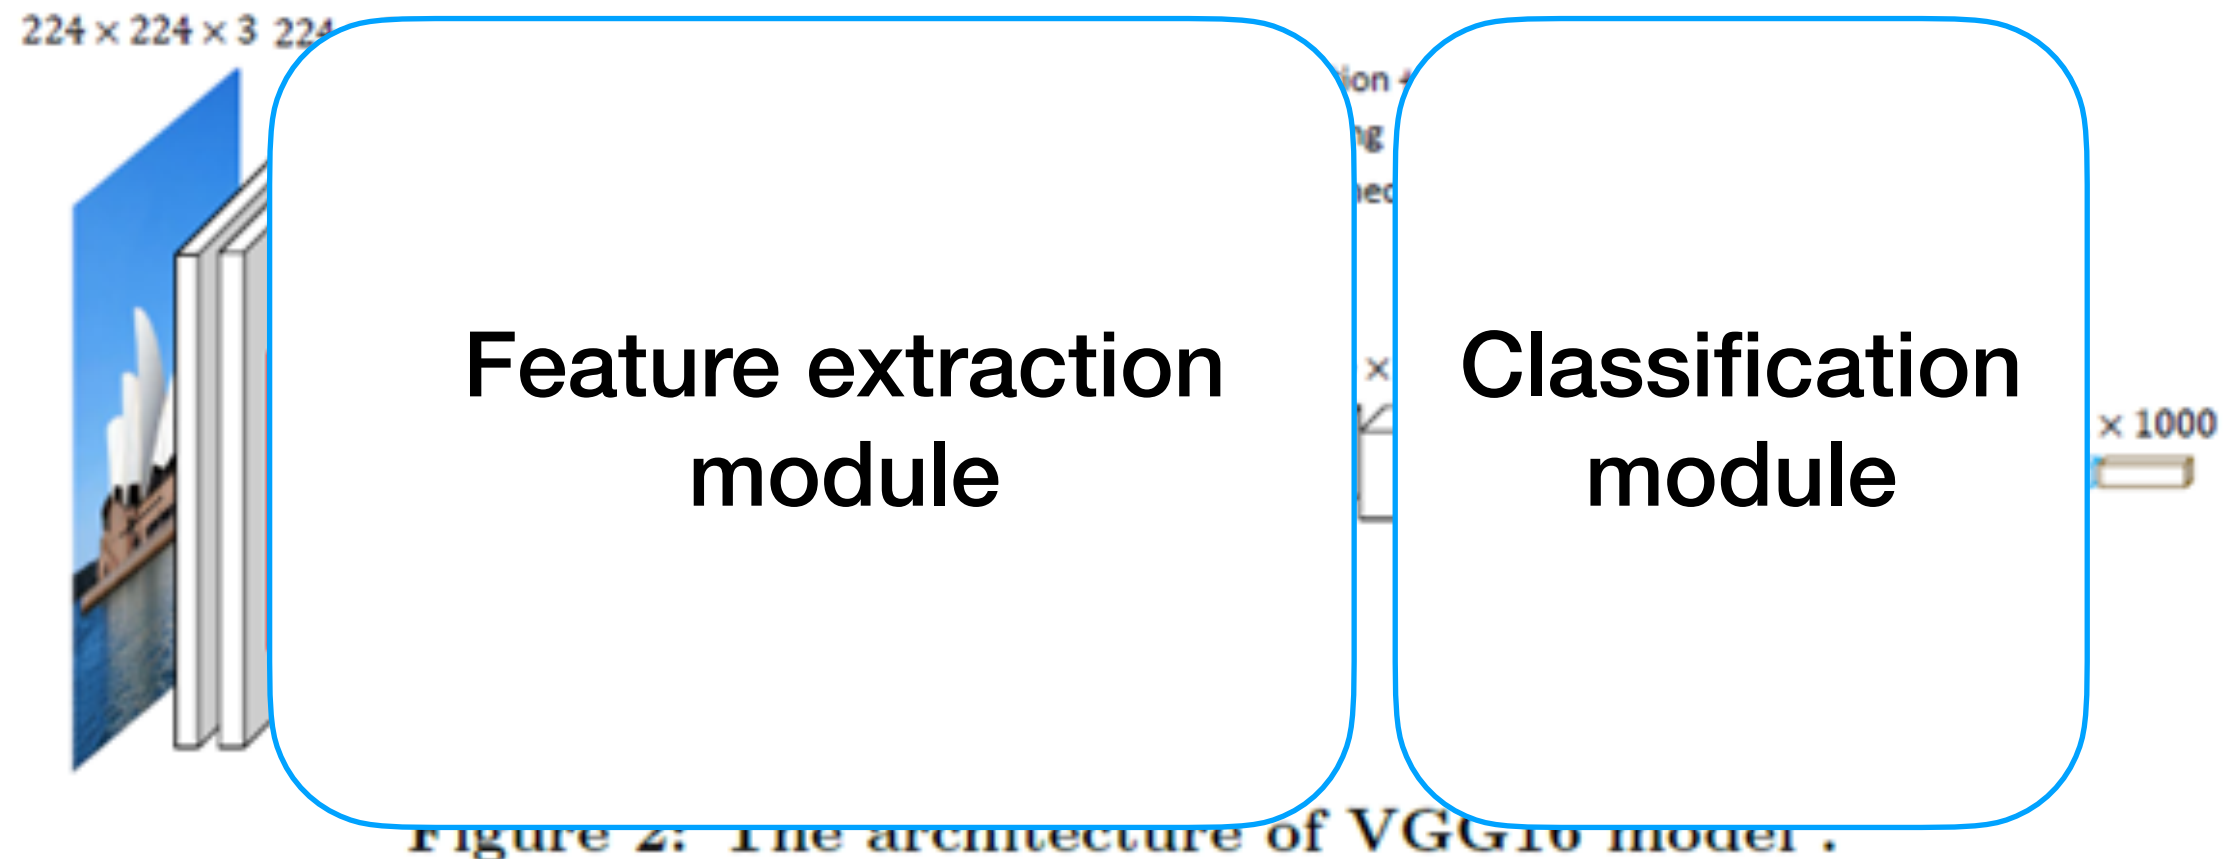


Figure 2: The architecture of VGG16 model .


General idea: the first layers involve generic feature extraction step and the last block can be regarded as a dataset-specific classification block.

Fine-tuning from pre-trained models




General idea: the first layers involve generic feature extraction step and the last block can be regarded as a dataset-specific classification block.

Fine-tuning from pre-trained models




https://github.com/CIA-Oceanix/DLCourse_MOi_2022/blob/main/notebooks/notebook_MNIST_classification_MLP_CNN_TransferLearning_students.ipynb

Over-fitting



Over-fitting



Regularization tricks to avoid overfitting

- Penalty terms In the training loss
- Data augmentation
- Dropout layers

Parameter norm penalization

- Regularized objective function:

$$\tilde{J}(\theta) = J(\theta) + \alpha\Omega(\theta)$$

- L² norm: $\Omega(\theta) = \frac{1}{2}||w||_2^2$
- L¹ norm: $\Omega(\theta) = ||w||_1 = \sum_i |w_i|$

Data augmentation

- Purpose: improving model generalization error by training on more data
- Very efficient for object recognition
- How to:
 - apply (geometric) transformations on input data (such as translation, rotation, scaling for images).
 - noise injection

Dropout





Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.


Examples of DL models for object recognition (2010-2020)



VGG16
(<100M of parameters)

Google Inception

(5M of parameters)



AlexNet
(60M of parameters)


Lecture. #2

Things to know (CNN)

- Convolution layers
- Pooling layers
- Activation layers
- Dropout layers
- Padding and stride
- Fine-tuning
- Over-fitting
- Data augmentation

Auto-encoders


Auto-encoders




Output with the same shape as the input

Application ?

Dense auto-encoders




PCA/EOF



PCA as a linear auto-encoder architecture.

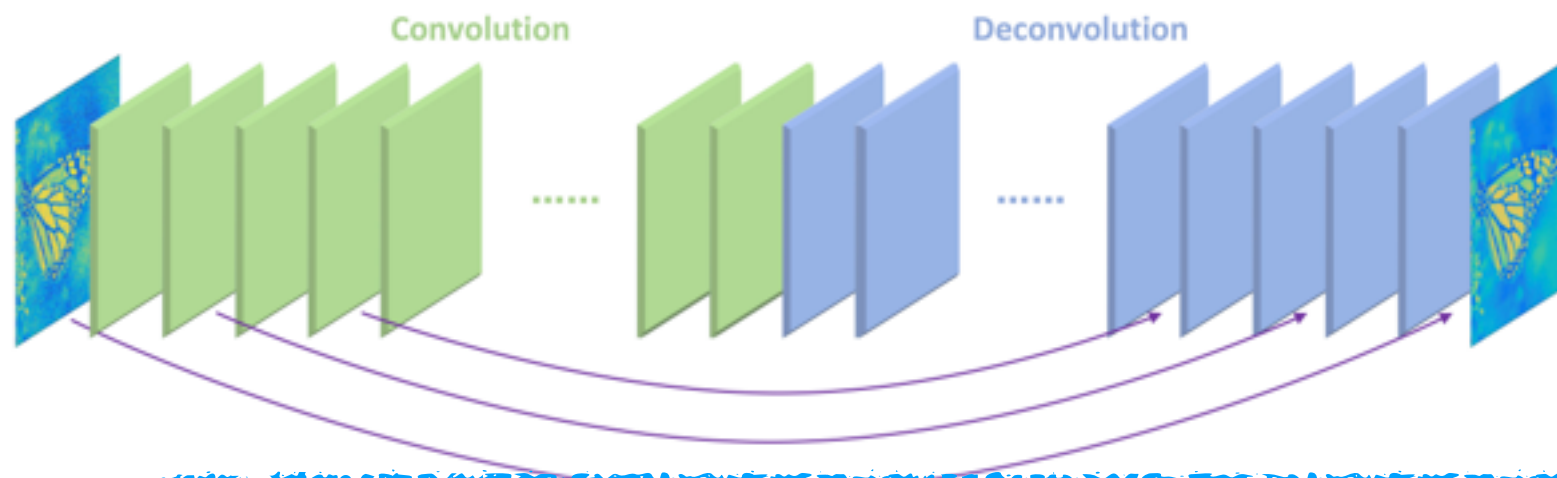
Which additional constraint ?

Convolutional auto-encoders




Convolutional AE Zoo

Many applications do not require a low-dimensional representation
(e.g., denoising, interpolation, super-resolution,...)




<https://arxiv.org/pdf/1606.08921.pdf>



UNet

Ronneberger
et al., 2015

Convolutional AE Zoo



Residual Block
(ResNet)

Rousseau et al.,
2019

Often used to address vanishing gradients (“very” deep networks)

Lecture. #2

Things to know (AE)

- Auto-encoder
- Latent variable
- UNet
- ResNet